

# Combined Dynamic Alpha + Selective N-Tokens Steering: Hallucination Guardrail Evaluation

**Summary:** This notebook evaluates a combined guardrail approach for Llama-3.1-8B that integrates both Dynamic Alpha (risk-proportional steering strength) and Selective N-Tokens Steering (applying the intervention only to the first N tokens). This method applies a dynamic, risk-scaled correction for high-risk prompts, but only during the initial generation steps, maximizing hallucination reduction while minimizing latency and preserving answer quality. This combined approach outperforms both individual ablations and is used for all further evaluations on other datasets.

- **Dynamic Alpha:** Steering strength (alpha) is scaled based on prompt risk, providing stronger correction for riskier prompts.
- **Selective N-Tokens:** Steering is applied only to the first 10 generated tokens, focusing intervention where it is most effective.

## Key Results (TruthfulQA Benchmark):

- **Baseline Model:** Accuracy: 38.57%, Hallucination Rate: 61.43%, Avg Latency: 3.86s
- **Combined Guarded Model:** Accuracy: 52.04%, Hallucination Rate: 47.96%, Avg Latency: 3.56s
- **Relative Error Reduction:** 21.93%
- **Latency Increase:** -7.78% (latency decreased)

This combined guardrail achieves the best trade-off between hallucination reduction, accuracy, and latency, and is therefore used for all subsequent cross-domain evaluations.

## Environment and Requirements Setup

Unzips the project folder and installs all required Python packages for Colab execution.

## Project Path and Colab Configuration

Adds the project directory to Python's path and sets the environment to 'colab' in the config file for compatibility.

```

import sys
import os

# Add project directory to Python's path
project_path = '/content/HallucinationVectorProject'
if project_path not in sys.path:
    sys.path.append(project_path)

# Programmatically set the environment to 'colab' in the config file
config_file_path = os.path.join(project_path, 'config.py')
with open(config_file_path, 'r') as f:
    lines = f.readlines()
with open(config_file_path, 'w') as f:
    for line in lines:
        if line.strip().startswith('ENVIRONMENT ='):
            f.write('ENVIRONMENT = "colab"\n')
        else:
            f.write(line)
print("Environment configured for Colab execution.")

```

Environment configured for Colab execution.

```

# Mount Google Drive
from google.colab import drive
drive.mount('/content/drive')

# Create a project directory to keep things organized
import os
PROJECT_DIR = "/content/drive/MyDrive/mistral-HallucinationVectorProj"
DATA_DIR = os.path.join(PROJECT_DIR, "data")
os.makedirs(DATA_DIR, exist_ok=True)

print(f"Project directory created at: {PROJECT_DIR}")

```

Start coding or generate with AI.

Start coding or generate with AI.

## Load Artifacts and Model Setup

Loads all required model artifacts, including the hallucination vector, risk classifier, and config thresholds, and prepares the model for evaluation.

```

import time
import pandas as pd
import torch
import csv
from tqdm import tqdm
import joblib

```

```
from unsloth import FastLanguageModel

# Import custom modules
import config
import utils

# This global dictionary will hold our models, tokenizer, vectors, etc.
artifacts = {}

def load_all_artifacts():
    """Loads all necessary model and project artifacts into the global
    if artifacts: return
    print("Loading all necessary artifacts for evaluation...")
    # Load the model from Hugging Face
    model, tokenizer = FastLanguageModel.from_pretrained(
        model_name = "unsloth/mistral-7b-instruct-v0.3-bnb-4bit",
        max_seq_length = max_seq_length,
        dtype = dtype,
        load_in_4bit = load_in_4bit,
    )

    model = FastLanguageModel.for_inference(model)
    model.gradient_checkpointing_disable()
    model.config.use_cache = True
    model.eval()

    artifacts['model'] = model
    artifacts['tokenizer'] = tokenizer
    artifacts['v_halluc'] = torch.load("/content/drive/MyDrive/mistral-7b-instruct-v0.3-bnb-4bit/v_halluc.pt")
    artifacts['risk_classifier'] = joblib.load("/content/Hallucination/risk_classifier.joblib")
    artifacts['thresholds'] = {
        "tau_low": config.TAU_LOW,
        "tau_high": config.TAU_HIGH,
        "optimal_alpha": config.OPTIMAL_ALPHA
    }

    # Load everything
load_all_artifacts()
```

```

Unsloth: Will patch your computer to enable 2x faster free finetuning
Unsloth Zoo will now patch everything to make training faster!
Loading all necessary artifacts for evaluation...
Loading model and tokenizer: unsloth/llama-3-8b-Instruct-bnb-4bit
==((=====))== Unsloth 2025.9.7: Fast Llama patching. Transformers: 4.5
    \\ /| Tesla T4. Num GPUs = 1. Max memory: 14.741 GB. Platform:
0^0/ \_/\ Torch: 2.8.0+cu126. CUDA: 7.5. CUDA Toolkit: 12.6. Triton
\       / Bfloat16 = FALSE. FA [Xformers = 0.0.32.post2. FA2 = Fal
"-_____" Free license: http://github.com/unslotha/unsloth
Unsloth: Fast downloading is enabled - ignore downloading bars which a
model.safetensors: 100% 5.70G/5.70G [00:40<00:00, 197MB/s]

generation_config.json: 100% 220/220 [00:00<00:00, 22.8kB/s]

tokenizer_config.json: 51.1k/? [00:00<00:00, 5.25MB/s]

tokenizer.json: 9.09M/? [00:00<00:00, 88.8MB/s]

special_tokens_map.json: 100% 345/345 [00:00<00:00, 24.5kB/s]

Model and tokenizer loaded successfully.

```

## Combined Selective Activation Steering and Guardrail Function

Defines a context manager to apply the steering vector only to the first N tokens, with dynamic risk-proportional strength, and a function to generate answers using this combined intervention.

```

import time
import torch
from contextlib import contextmanager

# Import our project's config and utils modules
import config
import utils

print("Defining combined logic for 'Dynamic Alpha + Selective Steering'")

# --- 1. The SelectiveActivationSteerer Class ---

class SelectiveActivationSteerer:
    def __init__(self, model, steering_vector, layer_idx, coeff=1.0,
                 self.model = model
                 self.vector = steering_vector
                 self.layer_idx = layer_idx
                 self.coeff = coeff
                 self.steering_token_limit = steering_token_limit
                 self._handle = None
                 self._layer_path = f"model.layers.{self.layer_idx}"
                 self.call_count = 0

    def _hook_fn(self, module, ins, out):
        self.call_count += 1
        if self.call_count <= self.steering_token_limit:

```

```
mistral-Dynamic_Alpha_&_Selective_N_Tokens_Ablation_project2_methodology2_hallucination_vector.ipynb - Colab

        steered_output = out[0] + (self.coeff * self.vector.to(out))
        return (steered_output,) + out[1:]
    return out

    def __enter__(self):
        self.call_count = 0
        try:
            layer = self.model.get_submodule(self._layer_path)
            self._handle = layer.register_forward_hook(self._hook_fn)
        except AttributeError:
            raise AttributeError(f"Could not find the layer at path: {self._layer_path}")
        return self

    def __exit__(self, exc_type, exc_val, exc_tb):
        if self._handle:
            self._handle.remove()

# --- 2. The New `answer_guarded_combined` Function ---
# This function combines the logic from both successful ablations.

def answer_guarded_combined(prompt_text: str, max_new_tokens: int = 1024):
    """Generates a response using the guardrail with DYNAMIC alpha and Selective N-Tokens Ablation.

    Args:
        prompt_text (str): The input prompt text.
        max_new_tokens (int): The maximum number of new tokens to generate. Default is 1024.

    Returns:
        str: The generated response.
    """
    start_time = time.time()

    risk_score = utils.get_hallucination_risk(
        prompt_text, artifacts['model'], artifacts['tokenizer'],
        artifacts['v_halluc'], artifacts['risk_classifier']
    )

    full_prompt = f"<|begin_of_text|><|start_header_id|>system<|end_header_id|>{prompt_text}"
    inputs = artifacts['tokenizer'](full_prompt, return_tensors="pt")
    input_token_length = inputs.input_ids.shape[1]

    if risk_score < artifacts['thresholds']['tau_high']:
        path = "Fast Path (Untouched)"
        with torch.no_grad():
            outputs = artifacts['model'].generate(**inputs, max_new_tokens=max_new_tokens)
    else:
        # From Dynamic Alpha Ablation: Calculate dynamic steering strategy
        optimal_alpha = artifacts['thresholds']['optimal_alpha']
        tau_high = artifacts['thresholds']['tau_high']
        scaling_factor = (risk_score - tau_high) / (1.0 - tau_high + optimal_alpha)
        dynamic_alpha = optimal_alpha * max(0, min(1, scaling_factor))

        path = f"Combined Steer Path (\u03b1={dynamic_alpha:.2f}, N={steering_token_limit})"

        # From Selective N-Tokens Ablation: Use the steerer with a to
        # We pass our newly calculated `dynamic_alpha` as the coefficient
        with SelectiveActivationSteerer(
            artifacts['model'], artifacts['v_halluc'], config.TARGET_TOKENS,
            coeff=dynamic_alpha,
            steering_token_limit=steering_token_limit
        ):
            outputs = artifacts['model'].generate(**inputs, max_new_tokens=max_new_tokens)

    return outputs[0].decode("utf-8")
```

```

        ):

            with torch.no_grad():
                outputs = artifacts['model'].generate(**inputs, max_n

            answer = artifacts['tokenizer'].decode(outputs[0], input_token_len
            latency = time.time() - start_time

        return {"answer": answer.strip(), "risk_score": risk_score, "path

    print("New function `answer_guarded_combined` is now defined and ready

Defining combined logic for 'Dynamic Alpha + Selective Steering' exper
New function `answer_guarded_combined` is now defined and ready for th

```

**Suppress Warnings** Suppresses specific sklearn warnings for cleaner output during evaluation.

```

import warnings

warnings.filterwarnings(
    "ignore",
    message="X does not have valid feature names",
    category=UserWarning,
    module="sklearn"
)

```

## Run Combined Guardrail Evaluation

Runs the evaluation loop on the TruthfulQA test set, applying the combined guardrail and saving results for each prompt.

```

# --- EXPERIMENT PARAMETER ---
STEERING_TOKEN_LIMIT = 10 # The 'N' for our selective steering

GUARDED_RESULTS_PATH_COMBINED = os.path.join("/content/HallucinationVectorProject", "guarded_results.csv")
BASELINE_RESULTS_PATH = os.path.join("/content/HallucinationVectorProject", "baseline_results.csv")

print(f"New guarded results will be saved to: {GUARDED_RESULTS_PATH_COMBINED}")

# Load the test set
test_df = pd.read_csv("/content/HallucinationVectorProject/data/final/test.csv")

# --- Resilient Evaluation Loop ---
guarded_headers = ['prompt', 'answer', 'risk_score', 'path_taken', 'latency']
utils.initialize_csv(GUARDED_RESULTS_PATH_COMBINED, guarded_headers)

processed_guarded = utils.load_processed_prompts(GUARDED_RESULTS_PATH_COMBINED)

print("Starting response generation for baseline and DYNAMIC guarded models")
for _, row in tqdm(test_df.iterrows(), total=len(test_df), desc="Dynamic Alpha + Selective N Tokens Ablation Project"):
    prompt = row['prompt']
    answer = row['answer']
    risk_score = row['risk_score']
    path_taken = row['path_taken']
    latency = row['latency']

    # Create inputs for DYNAMIC model
    dynamic_inputs = {
        "prompt": prompt,
        "max_new_tokens": STEERING_TOKEN_LIMIT,
        "temperature": 0.0,
        "top_p": 1.0,
        "top_k": 50,
        "do_sample": False,
        "logprobs": 0,
        "echo": False,
        "stop_sequences": []
    }

    # Create inputs for BASELINE model
    baseline_inputs = {
        "prompt": prompt,
        "max_length": 100,
        "temperature": 0.0,
        "top_p": 1.0,
        "top_k": 50,
        "do_sample": False,
        "logprobs": 0,
        "echo": False,
        "stop_sequences": []
    }

    # Generate responses
    dynamic_outputs = artifacts['model'].generate(**dynamic_inputs)
    baseline_outputs = artifacts['model'].generate(**baseline_inputs)

    # Decode and process outputs
    dynamic_answer = artifacts['tokenizer'].decode(dynamic_outputs[0], input_token_length=100)
    baseline_answer = artifacts['tokenizer'].decode(baseline_outputs[0], input_token_length=100)
    dynamic_risk_score = dynamic_outputs[1][0]
    baseline_risk_score = baseline_outputs[1][0]
    dynamic_path_taken = dynamic_outputs[2][0]
    baseline_path_taken = baseline_outputs[2][0]
    dynamic_latency = time.time() - start_time
    baseline_latency = time.time() - start_time

    # Save results
    result = {
        "prompt": prompt,
        "answer": dynamic_answer,
        "risk_score": dynamic_risk_score,
        "path_taken": dynamic_path_taken,
        "latency": dynamic_latency,
        "model": "DYNAMIC"
    }
    result["risk_score"] = dynamic_risk_score
    result["path_taken"] = dynamic_path_taken
    result["latency"] = dynamic_latency
    result["model"] = "DYNAMIC"

    result["risk_score"] = baseline_risk_score
    result["path_taken"] = baseline_path_taken
    result["latency"] = baseline_latency
    result["model"] = "BASELINE"

    # Write to CSV
    with open(GUARDED_RESULTS_PATH_COMBINED, "a") as f:
        writer = csv.DictWriter(f, fieldnames=guarded_headers)
        writer.writerow(result)

```

```

prompt = row['Question']

# Guarded Run
if prompt not in processed_guarded:
    try:
        result = answer_guarded_combined(prompt, steering_token_l
with open(GUARDED_RESULTS_PATH_COMBINED, 'a', newline='',
        csv.writer(f).writerow([prompt] + list(result.values))
    except Exception as e:
        print(f"Error on guarded prompt: {prompt}. Error: {e}")

# We already have Baseline Run results

print("Dynamic Alpha experiment generation complete.")

```

New guarded results will be saved to: /content/HallucinationVectorProject/GUARDED\_RESULTS\_PATH\_COMBINED  
Initialized CSV file at: /content/HallucinationVectorProject/results/combined.csv  
Starting response generation for baseline and DYNAMIC guarded models..  
Dynamic Alpha Evaluation: 100%|██████████| 617/617 [36:36<00:00, 3.56

## Run Judging, Analyze, and Summarize Results

Runs the judging process on generated answers, merges with ground truth, and computes final performance metrics for the combined guardrail experiment.

```

from evaluate_guardrail import run_judging_process
import utils
import config
import pandas as pd

# --- Redefine paths for the analysis ---
GUARDED_JUDGED_PATH_COMBINED = os.path.join("/content/HallucinationVectorProject/GUARDED_JUDGED_PATH_COMBINED")
BASELINE_JUDGED_RESULTS_PATH = os.path.join("/content/HallucinationVectorProject/BASELINE_JUDGED_RESULTS_PATH")
GUARDED_RESULTS_PATH_COMBINED = os.path.join("/content/HallucinationVectorProject/GUARDED_RESULTS_PATH_COMBINED")
BASELINE_RESULTS_PATH = os.path.join("/content/HallucinationVectorProject/BASELINE_RESULTS_PATH")

# Load the test set
test_df = pd.read_csv("/content/HallucinationVectorProject/data/final.csv")

# Load the newly generated results
guarded_df = pd.read_csv(GUARDED_RESULTS_PATH_COMBINED)
baseline_df = pd.read_csv(BASELINE_RESULTS_PATH)

# Merge with ground truth
guarded_merged_df = pd.merge(guarded_df, test_df, left_on='prompt', right_on='prompt')
baseline_merged_df = pd.merge(baseline_df, test_df, left_on='prompt', right_on='prompt')

# --- Run Judging ---
secrets = utils.load_secrets()
run_judging_process(guarded_merged_df, GUARDED_JUDGED_PATH_COMBINED,
# Assuming baseline is already judged, if not, uncomment below

```

```
# utils.run_judging_process(baseline_merged_df, config.BASELINE_JUDG

# --- Analyze and Print Final Report ---
guarded_judged_df = pd.read_csv(GUARDED_JUDGED_PATH_COMBINED)
baseline_judged_df = pd.read_csv(BASELINE_JUDGED_RESULTS_PATH)

baseline_accuracy = baseline_judged_df['is_correct'].mean()
guarded_accuracy = guarded_judged_df['is_correct'].mean()
baseline_error_rate = 1 - baseline_accuracy
guarded_error_rate = 1 - guarded_accuracy
relative_error_reduction = (baseline_error_rate - guarded_error_rate)
baseline_latency = baseline_judged_df['latency_seconds'].mean()
guarded_latency = guarded_judged_df['latency_seconds'].mean()
latency_increase_percent = (guarded_latency - baseline_latency) / ba

summary_data = {
    "Metric": ["Accuracy", "Hallucination Rate", "Avg Latency (s)",
    "Baseline Model": [f"{baseline_accuracy:.2%}", f"{baseline_error",
    "Guarded Model (Dynamic Alpha)": [f"{guarded_accuracy:.2%}", f"{
}
summary_df = pd.DataFrame(summary_data)

print("\n--- Final Performance Summary (Dynamic Alpha Experiment) --"
display(summary_df)
```

Loading secrets...

Secrets loaded successfully.

```
--- Starting Corrected Judging Process for combined_guarded_judged_results.csv
Initialized CSV file at: /content/HallucinationVectorProject/results/combo/
Found 0 already judged prompts. Resuming...
Judging combined_guarded_judged_results.csv: 100%|██████████| 617/617
--- Final Performance Summary (Dynamic Alpha Experiment) ---
```

	Metric	Baseline Model	Guarded Model (Dynamic Alpha)
0	Accuracy	38.57%	51.22%
1	Hallucination Rate	61.43%	48.78%
2	Avg Latency (s)	3.86	3.56
3	Relative Error Reduction	N/A	20.58%
4	Latency Increase	N/A	-7.80%